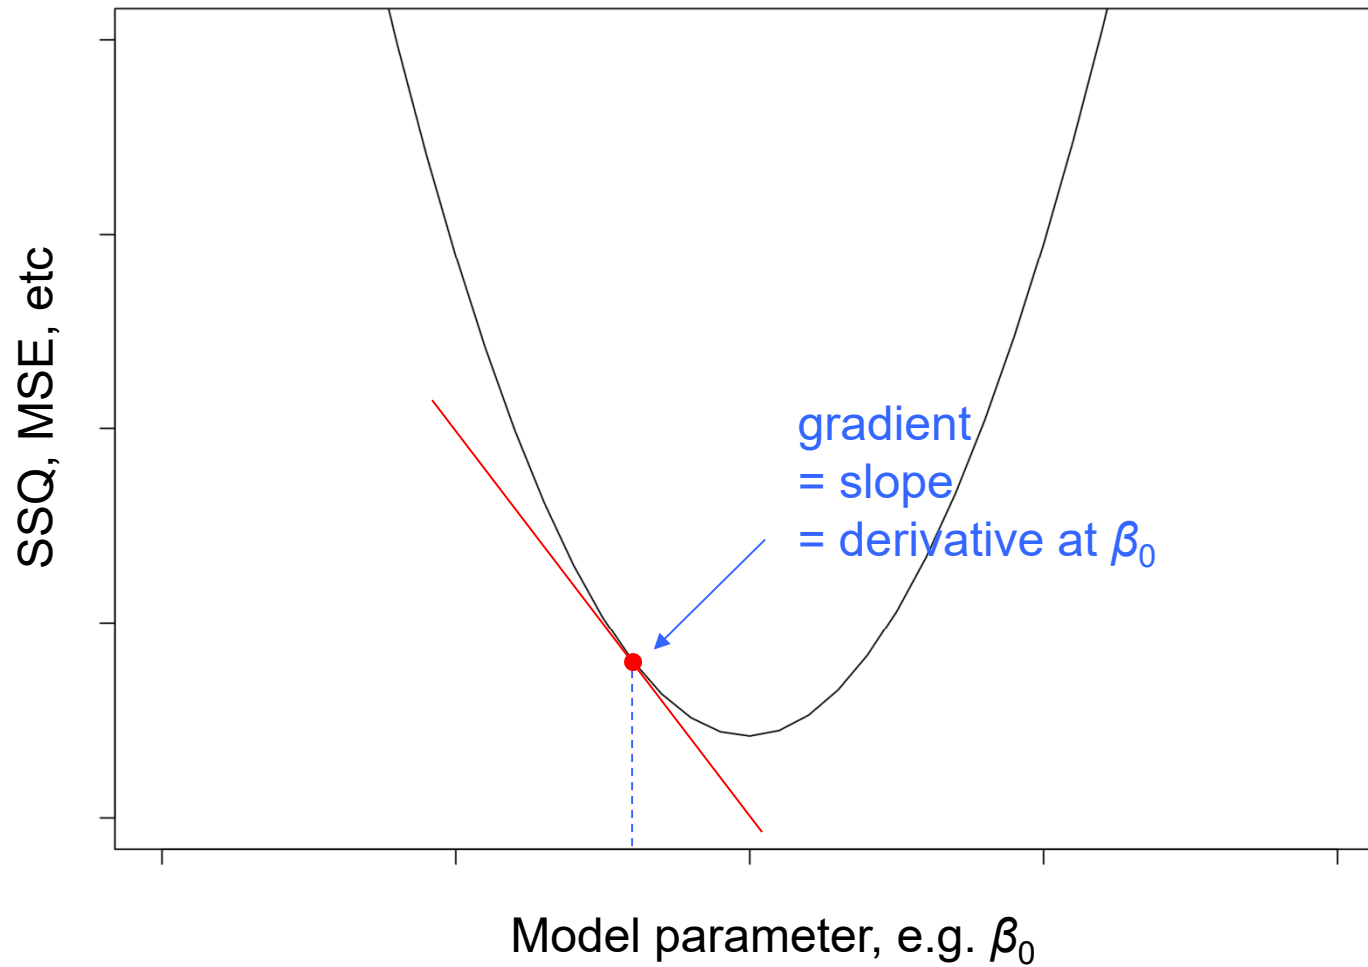


# Today

- Ensemble methods
  - Bagging
  - Random forest
  - Boosting
- But first
  - basic gradient descent
  - boosting is a variant

# Gradient descent



# Finding the gradient

- It turns out that the gradient for a linear model is a function of the residuals
- See math next

# Finding the gradient

Loss  
function

$$SSQ = \sum_i^n (y_i - \hat{y}_i)^2$$

$$= \sum_i^n r_i^2$$

$$= \sum_i^n (y_i - \beta_0 - \beta_1 x_i)^2$$

$$= \sum_i^n y_i^2 - 2y_i\beta_0 - 2y_i\beta_1 x_i + \beta_0^2 + 2\beta_0\beta_1 x_i + \beta_1^2 x_i^2$$

$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

$$r_i = y_i - \beta_0 - \beta_1 x_i$$

first derivatives

$$\frac{\partial SSQ}{\partial \beta_1} = \sum_i^n (-2y_i x_i + 2\beta_0 x_i + 2\beta_1 x_i^2)$$

$$= \sum_i^n -2x_i (y_i - \beta_0 - \beta_1 x_i)$$

$$= -2 \sum_i^n r_i x_i$$

x is constant

gradient proportional to r

Gradient =  
first derivative

# Finding the gradient

Loss  
function

$$SSQ = \sum_i^n (y_i - \hat{y}_i)^2$$

$$= \sum_i^n r_i^2$$

$$= \sum_i^n (y_i - \beta_0 - \beta_1 x_i)^2$$

$$= \sum_i^n y_i^2 - 2y_i\beta_0 - 2y_i\beta_1 x_i + \beta_0^2 + 2\beta_0\beta_1 x_i + \beta_1^2 x_i^2$$

$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

$$r_i = y_i - \beta_0 - \beta_1 x_i$$

first derivatives

$$\frac{\partial SSQ}{\partial \beta_0} = \sum_i^n -2y_i + 2\beta_0 + 2\beta_1 x_i$$

Gradient =  
first derivative

$$= -2 \sum_i^n r_i \rightarrow \text{gradient proportional to } r$$

# Gradient descent

Gradient descent training algorithm for a linear model

```
set lambda (learning rate)
make initial guess for  $\beta_0, \beta_1$ 
for many iterations
    find gradient at  $\beta_0, \beta_1$ 
    step down:  $\beta = \beta - \text{lambda} * \text{gradient}(\beta)$ 
print  $\beta_0, \beta_1$ 
```

# Gradient descent

- Code
  - `gradient_descent.R`
  - gradient descent for linear model

# Gradient boosting

## Gradient boosting algorithm (intuitive version)

set  $\lambda$  (learning rate)

fit a model,  $m(x)$ , to the data

keep a fraction of the model,  $\lambda m(x)$

calculate the left over variation,  $r$  (not explained by the model fraction)

repeat until no more systematic variation in  $r$

    fit a model,  $m(x)$ , to  $r$

    keep a fraction of the model,  $\lambda m(x)$

    add to previous fraction

    calculate the left over variation,  $r$  (not explained by summed fractions)

prediction is the final sum of model fractions



# Gradient boosting

## Algorithm

load  $y, x, x_{\text{new}}$

Set hyper-parameters:  $\lambda, \dots$

Can be any  
model



set  $\hat{f}(x_{\text{new}}) = 0$

set  $r \leftarrow y$  (residuals equal to the data)

for  $m$  in 1 to  $n_{\text{iterations}}$

train model on  $r$  and  $x$

predict residuals,  $\hat{r}_m(x)$ , from trained model

update residuals:  $r \leftarrow r - \lambda \hat{r}_m(x)$

Gradient  
descent



predict  $y$  increment,  $\hat{f}_m(x_{\text{new}})$ , from trained model

update prediction:  $\hat{f}(x_{\text{new}}) \leftarrow \hat{f}(x_{\text{new}}) + \lambda \hat{f}_m(x_{\text{new}})$

return  $\hat{f}(x_{\text{new}})$

# Gradient descent

predict residuals,  $\hat{r}_m(x)$ , from trained model

update residuals:  $r \leftarrow r - \lambda \hat{r}_m(x)$


**Loss function** is MSE and we are descending its surface

**Gradient** at  $x$  is proportional to the predicted residual  $\hat{r}_m(x)$

$\lambda \hat{r}_m(x)$  is the **increment** taken down the gradient

$r$  gets closer to 0 at each step, so MSE goes down

Boosted linear model (black line) compared to linear regression (red dashed line). They are the same.



# Gradient descent

- Code
  - `gradient_descent.R`
  - gradient boosting for linear model

# Boosted regression tree

## Algorithm

load  $y, x, x_{\text{new}}$

set parameters: `tree_complexity`, `ntrees`,  $\lambda$

set  $\hat{f}(x_{\text{new}}) = 0$

set  $r \leftarrow y$  (residuals equal to the data)

for  $m$  in 1 to `ntrees`

    train `tree model`,  $m$ , on  $r$  and  $x$

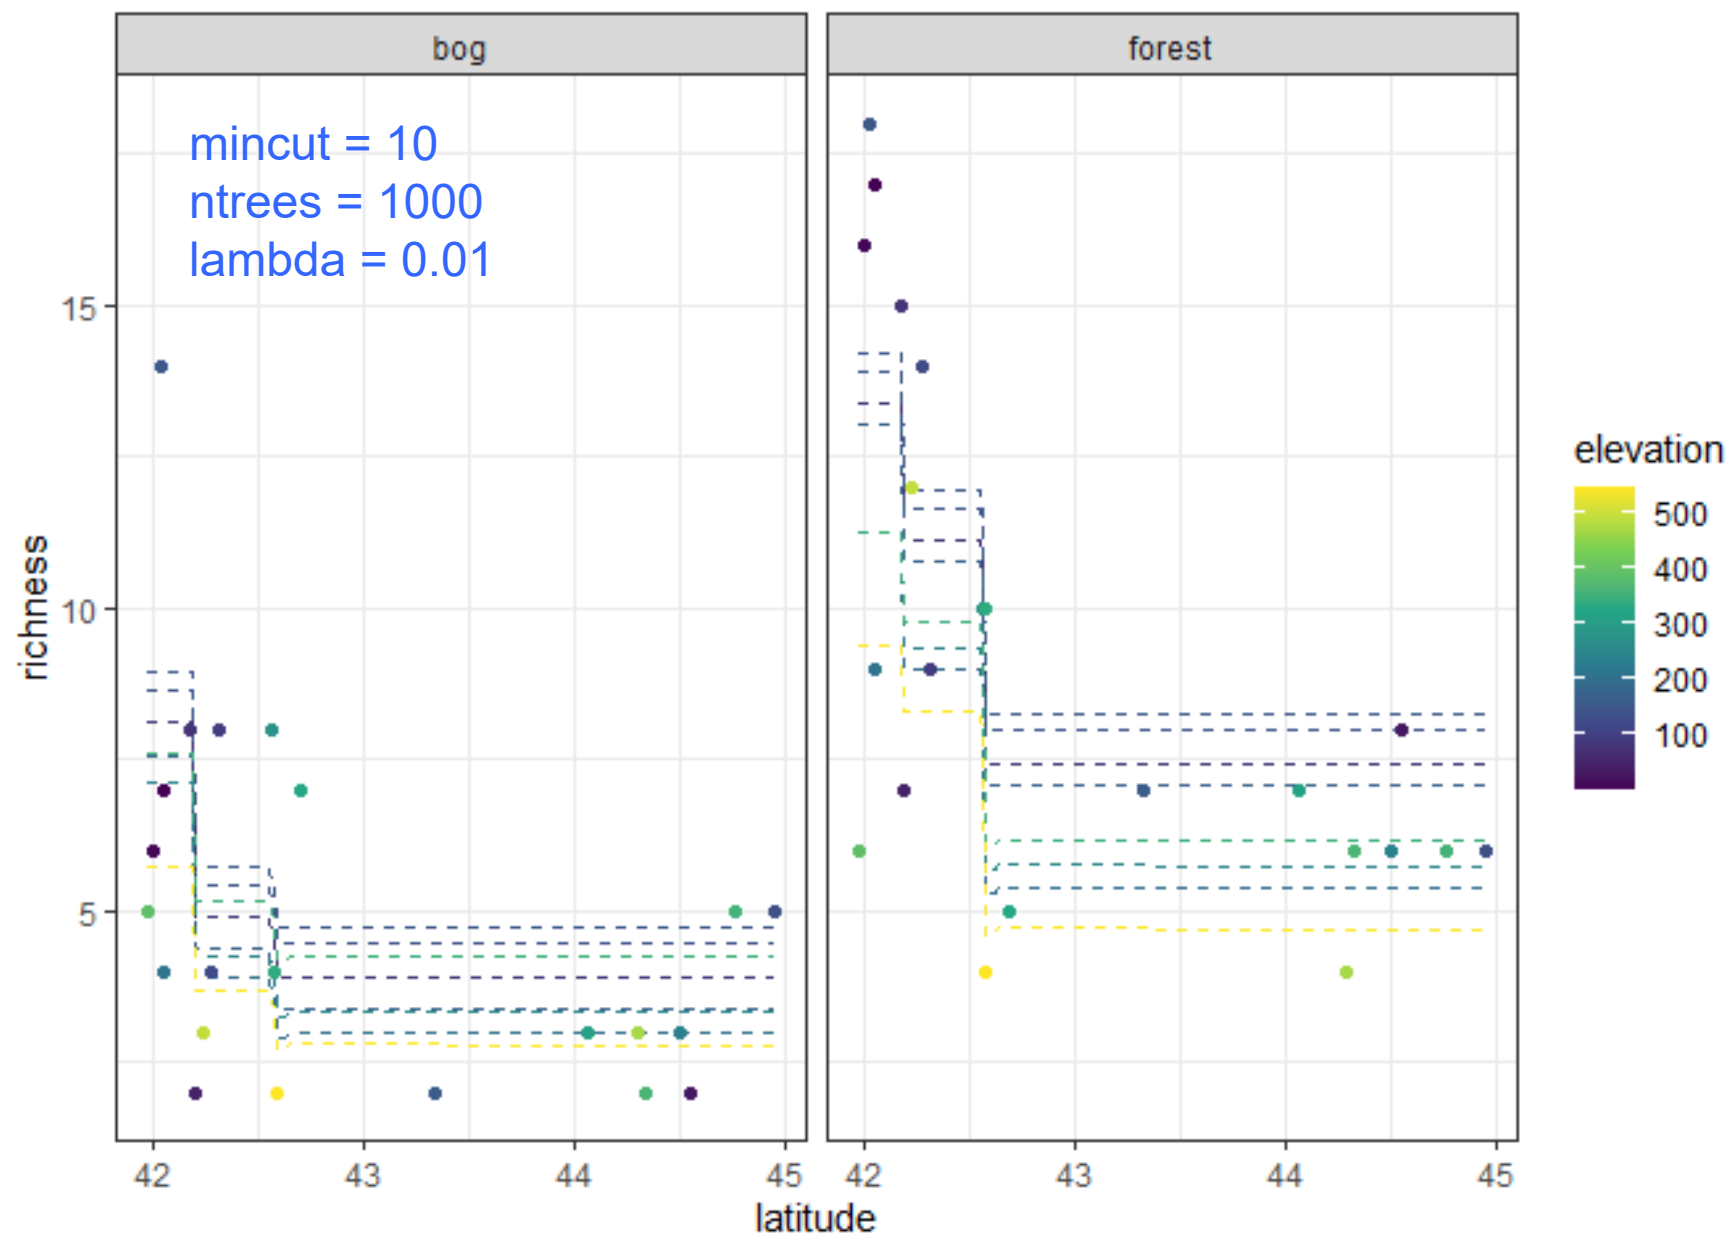
    predict residuals,  $\hat{r}_m(x)$ , from `trained tree`

    update residuals:  $r \leftarrow r - \lambda \hat{r}_m(x)$

    predict  $y$  increment,  $\hat{f}_m(x_{\text{new}})$ , from `trained tree`

    update prediction:  $\hat{f}(x_{\text{new}}) \leftarrow \hat{f}(x_{\text{new}}) + \lambda \hat{f}_m(x_{\text{new}})$

return  $\hat{f}(x_{\text{new}})$



# Gradient descent

